

A Clinical Trials ChatBot that Speaks Your Language

Douglas Bailey, University of Michigan, UMSI

Dmitry Danilov, University of Michigan, UMSI

Parker Moesta, University of Michigan, UMSI

“Information is cheap, but meaning is expensive.” — George Dyson

INTRODUCTION

Confronted by a troubling medical issue, individuals turn to the internet seeking information and hope from the latest clinical trials. Yet what they often find is tortuous technical terminology, carefully crafted commercials from big-Pharma, and manipulated misinterpretations. Even the latest Large Language Models (LLMs) can come up short with hallucinated responses (Figure 1).

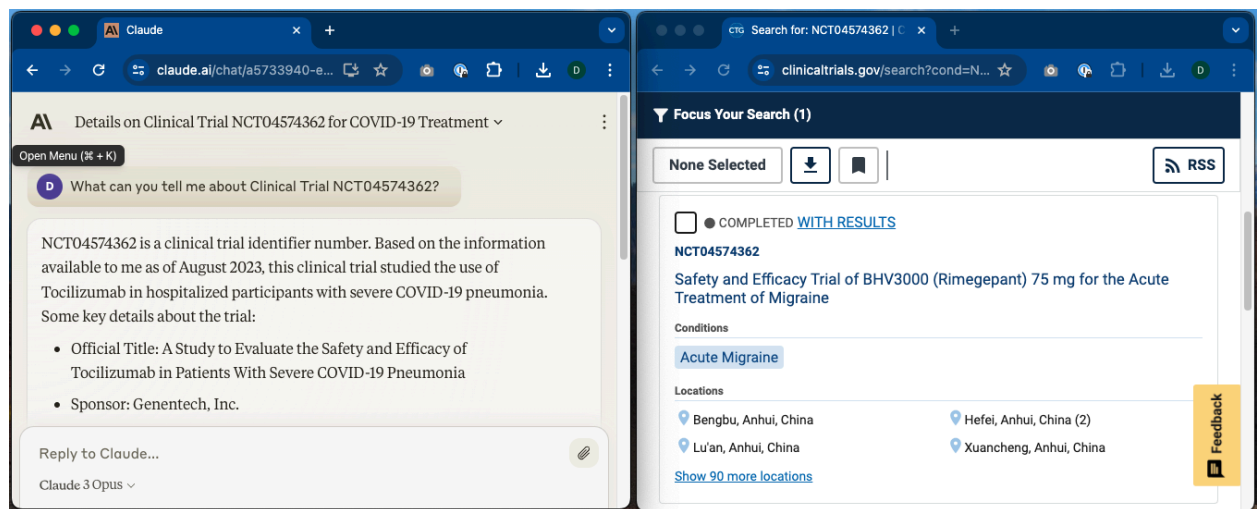


Figure 1. Large Language Models (LLM): *The Latest and Greatest?*

Left: Anthropic Claude 3 Opus (accessed 14 April 2024) returns a spurious response.

Right: The “ground truth” response from ClinicalTrials.gov

Our project takes a step towards solving this problem by providing a ChatBot that interfaces directly with data from ClinicalTrials.gov. Provided by the National Library of Medicine (NLM), ClinicalTrials.gov is an authoritative, trusted database with over 430,000 studies.

Based on Retrieval Augmented Generation (RAG) technology (Lewis et al., 2021), our ChatBot allows users to interact and receive high-quality, plain language information derived from this reliable source. This approach allows users to extract meaning from complex information.

PROJECT OVERVIEW

Clinical Trials — Background

Clinical trials are classified as interventional or observational (American Cancer Society, n.d.). Randomized controlled trials (RCTs) of treatments (e.g., medicines) versus control (e.g., a placebo) are interventional. As shown in Figure 2, within this group:

- Phase 1 studies explore tolerable dosage.
- Phase 2 studies evaluate treatment efficacy.
- Phase 3 studies often use RCTs to confirm efficacy and safety.
- Phase 4 deals with approved treatments.




Figure 2. Clinical Trials Phases 1 through 4

(Note: not exhaustive, classifications such as “Early Phase1 and Phase0 exist, but are beyond the scope of this paper.)




Figure 3. It's complicated — NetworkX Representation of Clinical Trial NCT01841944 —
Files can contain data spanning thousands of lines.

A (Mercifully) Brief Introduction to the Retrieval Augmented Generation (RAG)

To get an intuitive sense of how Retrieval Augmented Generation (RAG) works, we can unpack the acronym in reverse order. "Generation" refers to the natural language output produced by a Large Language Model (LLM), sometimes oversimplified as next token generation. Seeking to refine the generated output, a user can employ prompt engineering to "augment" the context the LLM uses as the basis for its output. By systematically "retrieving" content from a known source and supplying that information to the LLM, we complete the RAG workflow — a process of retrieval and augmentation to enhance the quality of the generated text. As a result, RAG generates coherent and relevant output (Figure 4).

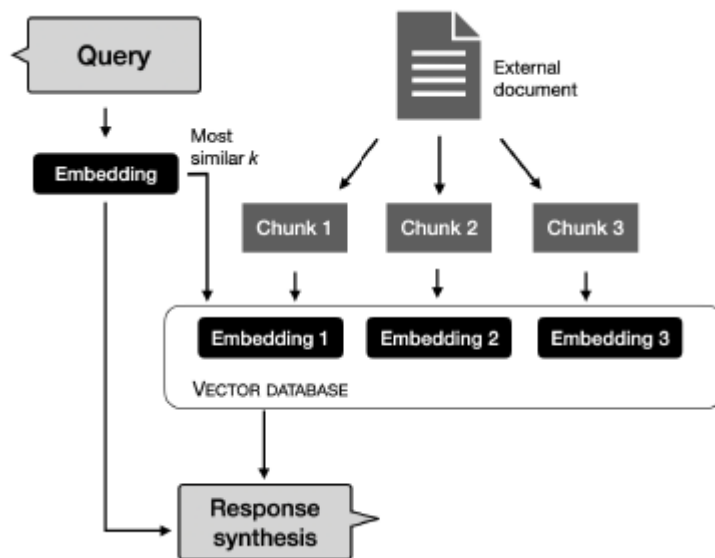


Figure 4. Schematic Overview of a "Naive" RAG (Raschka, 2023)

Applying RAG to clinicaltrials.gov

Advantages

- Uses LLM in inference mode, not for data retrieval—less hallucination, misinformation)
- Source of information (nodes) is observable and explainable—unlike inherently inscrutable LLMs)
- Can update with the latest information — not subject to LLMs' "knowledge cutoff" date
- In contrast to a response from a SQL query, RAG outputs information as user-friendly text

Project Objectives

- Implement a simple user interface for a RAG application on ClinicalTrials.gov data
- Allow users to query a clinical trial to discover basic facts
- Provide a Plain Language Summary (e.g., can explain statistical metrics at a high-school level)

Related Work

Zakka et al. (2023) in "Almanac: Retrieval-Augmented Language Models for Clinical Medicine" write that "large language models [are] effective tools in the clinical decision-making process." While the dataset and approach differ from ours, the paper's discussion of evaluation is recommended.

Huang (2023) discusses a similar objective in the Medium article "Build a Chatbot for Clinical Trials Across Multiple Data Sources," although Huang's approach does not use retrieval-augmented generation (RAG).

Previous work sought to facilitate access to the clinicaltrials.gov database. We take advantage of one such effort: the Clinical Trials Transformation Initiative (AAct). The AAct is a collaboration between Duke University and the U.S. Food and Drug Administration (FDA) (Clinical Trials Transformation Initiative, n.d.). It provides a PostgreSQL database corresponding exactly to clinicaltrials.gov.

DATA SOURCES AND PROPERTIES

Our project relies on public domain data sources with no usage restrictions.

ClinicalTrials.gov is a comprehensive database provided by the National Library of Medicine (NLM). It includes both interventional and observational studies for clinical trials run in over 200 countries. Each trial is subject to the NLM review process.

In the initial phase of our project, we focused on a subset of the vast number of available clinical trials. This subset specified: *Phase 3, Study Completed, Sponsored by Pfizer*, and with *Reported p-values*, comprising approximately 250 trials.

During testing we scaled-up our dataset by an order of magnitude to approximately 5,600 trials. This larger dataset includes trials with a variety of sponsors and more diverse data. As we scaled-up, we were pleased to find that our data ingestion, cleaning, and RAG processing workflows remained effective without needing further modification. This gives us confidence that we could handle even larger datasets in a future extension of the project.

SYSTEM DESIGN

The main components of our system as well as the overall end-to-end workflow are outlined in Figure 5.

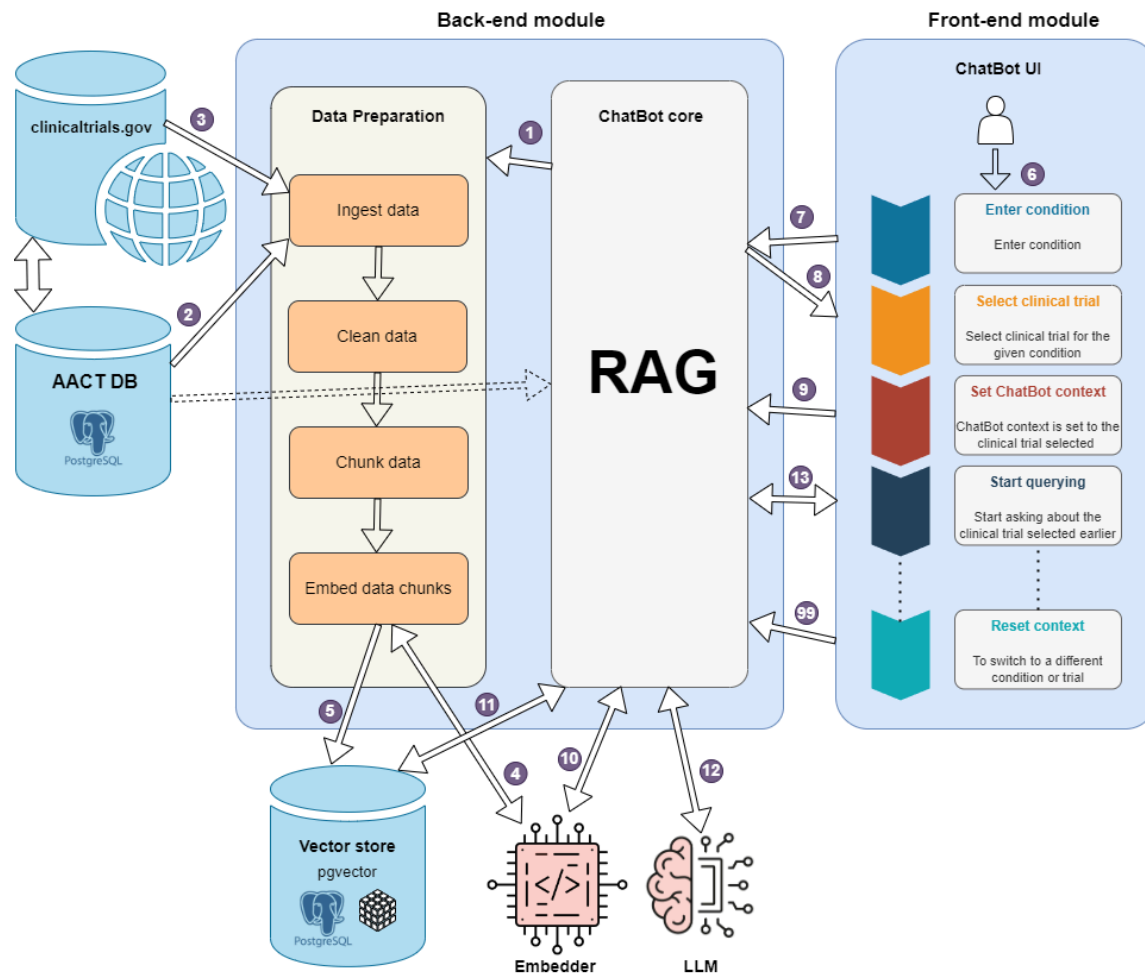


Figure 5. Main components of the system and end-to-end workflow

Workflow stages (marked as circled numbers in the diagram):

1. The system initiates data ingestion.
2. IDs of the clinical trials are extracted from the **AACT database**.
3. Clinical trials are downloaded from **clinicaltrials.gov** using their public API.
4. Clinical trials data are cleaned and converted into documents, documents are then split into chunks.
5. For each chunk, a fixed-size vector of embeddings is calculated using an embedding **LLM**.
6. Vectors of embeddings are stored in the vector database (**Postgres + pgvector** extension).
7. User initiates the chatbot and enters a condition in the search field.
8. Using SQL, system returns a list of clinical trials associated with the condition entered.
9. User selected a clinical trial he/she is interested in learning more about — this sets the global chat session context.
10. User enters a query to ask about the clinical trial selected.
11. System converts the query into a vector of embeddings and uses this vector to retrieve chunks of documents that are semantically similar to the query from the vector database.

12. Both the query and retrieved chunks of semantically similar documents are added to the synthesizer **LLM**'s context. The **LLM** uses the context to synthesize a response to the user's query. Response is displayed on the UI.
13. Steps 10-12 are repeated for as long as the user has more questions to ask
...
14. Once the user is done querying the clinical trial, he/she resets the chat session context. The user is redirected to the start page and can start the workflow over again using different conditions and/or clinical trials as needed.

DISCUSSION

Unlike evaluating supervised Machine Learning, evaluating the performance of a RAG system is challenging; mainly because there are no ground-truth labels that correspond exactly with model output. Therefore, metrics such as BLEU scores and ROUGE scores, which guide researchers in improving LLM output, are not entirely adequate. For example, a BLEU score uses n-gram overlap to determine how similar strings are, but it doesn't always provide an accurate reflection of meaning. For instance, strings such as "*Her favorite animal is a cat with long whiskers*" and "*Her favorite animal is a catfish with long whiskers*" have a high degree of n-gram overlap, yet the meaning is different.

During the course of our project, we found ourselves interacting with RAG on an ad hoc basis, making observations and adjusting model parameters as we developed a feel for the system's capabilities and limitations. As others have noted, "human evaluations are tedious, expensive, hard-to-automate, and subjective" (Raschka, 2023). Accordingly, we appreciated the "LLM-as-a-judge" approach advanced in "*Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena*" (Zheng et al., 2023). Because our RAG system is based on libraries and dependencies provided by LlamaIndex, we took advantage of some pre-built metrics and example code to develop a metrics framework (LlamaIndex Documentation, n.d.).

Specifically, we experimented with:

1. "Faithfulness" — how well the response matches the retrievable context — here, a very low score indicates hallucination.
2. "Relevancy" — measures whether the response and context match up nicely with the query. A low score suggests the answer might be off-topic or not addressing the specific question (like students who answer what they know, rather than what the teacher asked).
3. "Hit Rate" — the percentage of times the most relevant document falls within the top-k retrieved documents by the RAG system. This helps gauge the capability of the retrieval.
4. "Mean Reciprocal Rank" (MRR) — a perfect MRR score of 1 indicates the most relevant document is always retrieved first. But if the RAG returns the most relevant document as its 3rd choice, that would be $MRR=0.33$ (thus, higher is better).

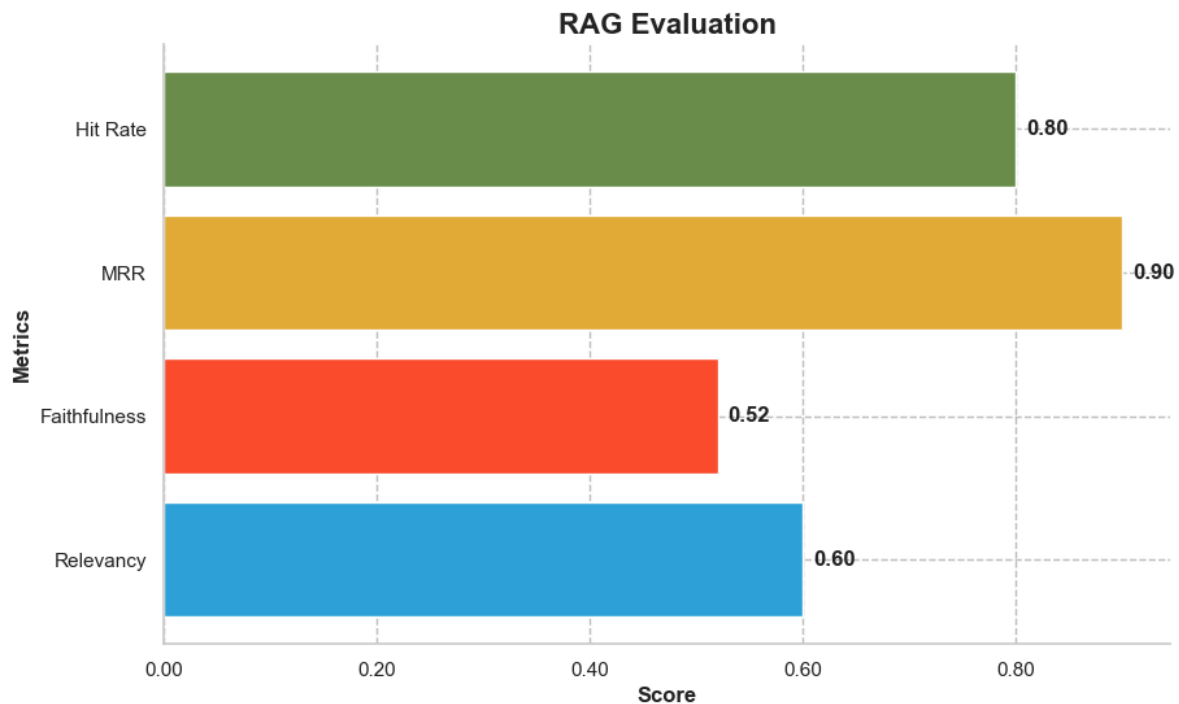


Figure 6. RAG Evaluation Metrics

Ethical Considerations

We recognize that our solution must not oversimplify research findings to the extent that it returns distorted or erroneous information. We don't want users to base critical decisions on information that may be well-intentioned, but flawed.

Mitigation Strategies

Accordingly, users must be properly informed about the purpose of our project, which is simply to provide an ancillary source of information. In the project's current iteration users will encounter a standard disclaimer declaring: "information only, not medical advice." Perhaps a future project iteration will have a more interactive feature, such as a checkbox or a prompt that informs and confirms the user's understanding of system limitations. We may also consider taking advantage of RAG ability to return links to retrieved context information — this way original source information can be easily referenced. Ultimately, we could employ "focus group" methods to study user interaction with the system. Qualified researchers could assist in the evaluation of user experience.

CONCLUSION

Giving people an easy-to-use interface to get comprehensible information from ClinicalTrials.gov is the project objective. We believe providing plain language summaries of research information to those who seek it, but lack scientific sophistication, has genuine value.

This project explored using Retrieval Augmented Generation (RAG) as the basis for a ChatBot interface with ClinicalTrials.gov. The system is able to answer basic questions about clinical trials and provide plain language summaries to help users interpret complex trial information.

The main achievements of the project include;

- Ingesting and cleaning a dataset of over 5,000 heterogeneous clinical trials
- Applying the RAG framework

- Implementing a simple ChatBot
- Experimenting with evaluation metrics for retrieval and generation quality

Limitations of the current implementation are:

- Dataset still represents a fraction of all available trials on ClinicalTrials.gov
- Evaluation of RAG outputs remains largely subjective and ad hoc
- Potential for misinterpretation of simplified trial information by lay users

Potential next steps:

- Scaling up the dataset by an order of magnitude
- Continued evaluation of RAG metrics
- Enhancing the ChatBot interface and
- Implementing a better disclaimers to mitigate risks of misuse
- Returning links to source information for reference and validation

Our project aspiration was to take a step towards making complex clinical trial information more accessible and interpretable for a general audience. Our task wasn't easy; we needed to develop new skills and overcome limitations inherent in our application of Retrieval Augmented Generation. Our experience suggests that future developments could serve as a valuable public resource.

STATEMENT OF WORK

This project team's collaboration was inspired by the pair programming of Google's Jeff Dean and Sanjay Ghemawat (Somers, 2018). We emulated this approach in our project by working together towards a common goal and maximizing each member's unique skills.

REFERENCES

American Cancer Society. (n.d.). Phases of clinical trials. Retrieved April 15, 2024, from <https://www.cancer.org/cancer/managing-cancer/making-treatment-decisions/clinical-trials/what-you-need-to-know/phases-of-clinical-trials.html>

Clinical Trials Transformation Initiative. (n.d.). AACT Database. Retrieved April 16, 2024, from <https://aact.ctti-clinicaltrials.org/>

Huang, S. (2023, July 29). Build a Chatbot for Clinical Trials Across Multiple Data Sources. Medium.

<https://medium.com/star-gazers/build-a-chatbot-for-clinical-trials-across-multiple-data-sources-f121211cec98>

Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-t., Rocktäschel, T., Riedel, S., & Kiela, D. (2021). Retrieval-augmented generation for knowledge-intensive NLP tasks. arXiv:2005.11401v4 [cs.CL]. <https://doi.org/10.48550/arXiv.2005.11401>

LlamaIndex Documentation. (n.d.). Retrieval evaluation. Retrieved from https://docs.llamaindex.ai/en/stable/examples/evaluation/retrieval/retriever_eval/

Raschka, S. (2023). Machine learning Q and A: Expand your machine learning & AI knowledge with 30 in-depth questions and answers. Leanpub. <http://leanpub.com/machine-learning>

Somers, J. (2018, December 3). The friendship that made Google huge. The New Yorker. Retrieved from <https://www.newyorker.com/magazine/2018/12/10/the-friendship-that-made-google-huge>

Zakka, C., Chaurasia, A., Shad, R., Dalal, A. R., Kim, J. L., Moor, M., Alexander, K., Ashley, E., Boyd, J., Boyd, K., Hirsch, K., Langlotz, C., Nelson, J., & Hiesinger, W. (2023). Almanac: Retrieval-Augmented Language Models for Clinical Medicine. Research Square. <https://doi.org/10.21203/rs.3.rs-2883198/v1>

Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., & Stoica, I. (2023). Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. arXiv. <https://doi.org/10.48550/arXiv.2306.05685>